New York City Taxi Trip Duration

Submitted to:
Prof. Douglas Jones

Group Members:

Divya Naidu
Hemakshi Shardha
Shwetha Krishna
Tarun Singh

December 03, 2018

Business Problem:

Build a model that predicts the total ride duration of taxi trips in New York City

Overview

Source: https://www.kaggle.com

Data Fields:

- Id
- Vendor_id
- Pickup_datetime
- Dropoff_datetime
- Passenger_count
- Pickup_longitude
- Pickup_latitude
- Dropoff_longitude
- Dropoff_latitude
- Store_and_fwd_flag
- Trip_duration

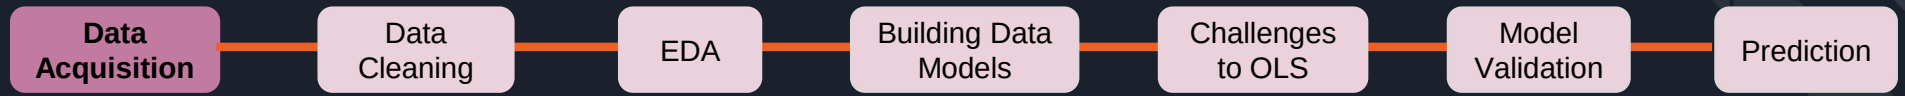

Libraries Used

- library(dply)
- library(tibble)
- library(tidyr)
- library(stringr)
- library(forcast)
- library(lubridate)
- library(Amelia)
- library(mice)
- library(moments)
- library(ggplot2)
- library(rgdal)
- library(data.table)
- library(dplyr)
- library(geosphere)
- library(car)
- library(corrplot)
- library(DAAG)
- library(faraway)
- library(GGally)
- library(corrplot)
- library(gridExtra)

Data Manipulation

- Converted 0's to NAs
- Converted categorical variable into a factor variable
- Converted the format of Data and time variables
- Used missmap to identify missing data.
- Used md.pattern to check the pattern of missing data
- Used MICE package for multiple imputation

Summary- Before MICE

From summary, we can see that there are NAs in following columns:

- Pickup_Longitude
- Pickup_Latitude
- Dropoff_Longitude
- Dropoff_Latitude
- Trip_Duration

```
     id            vendor_id pickup_datetime
Length:1500        1:719     Min.   :2016-01-01 00:20:00
Class :character   2:781     1st Qu.:2016-02-19 19:27:45
Mode  :character             Median :2016-04-03 20:27:30
                             Mean   :2016-04-02 15:30:06
                             3rd Qu.:2016-05-14 20:18:15
                             Max.   :2016-06-30 23:21:00

dropoff_datetime              passenger_count pickup_longitude
Min.   :2016-01-01 00:24:00   1:1044          Min.   :-74.19
1st Qu.:2016-02-19 19:45:00   2: 239          1st Qu.:-73.99
Median :2016-04-03 20:43:00   3:  64          Median :-73.98
Mean   :2016-04-02 15:44:59   4:  32          Mean   :-73.97
3rd Qu.:2016-05-14 20:41:15   5:  80          3rd Qu.:-73.97
Max.   :2016-06-30 23:37:00   6:  41          Max.   :-73.78
                                              NA's   :166
pickup_latitude dropoff_longitude dropoff_latitude store_and_fwd_flag
Min.   :40.63   Min.   :-74.18    Min.   :40.59    N:1492
1st Qu.:40.74   1st Qu.:-73.99    1st Qu.:40.73    Y:   8
Median :40.75   Median :-73.98    Median :40.76
Mean   :40.75   Mean   :-73.97    Mean   :40.75
3rd Qu.:40.77   3rd Qu.:-73.96    3rd Qu.:40.77
Max.   :40.88   Max.   :-73.78    Max.   :41.00
NA's   :179     NA's   :158       NA's   :144
trip_duration
Min.   :   78.0
1st Qu.:  431.0
Median :  682.5
Mean   :  904.8
3rd Qu.: 1076.2
Max.   :86137.0
NA's   :20
```

Missmap from Amelia Library

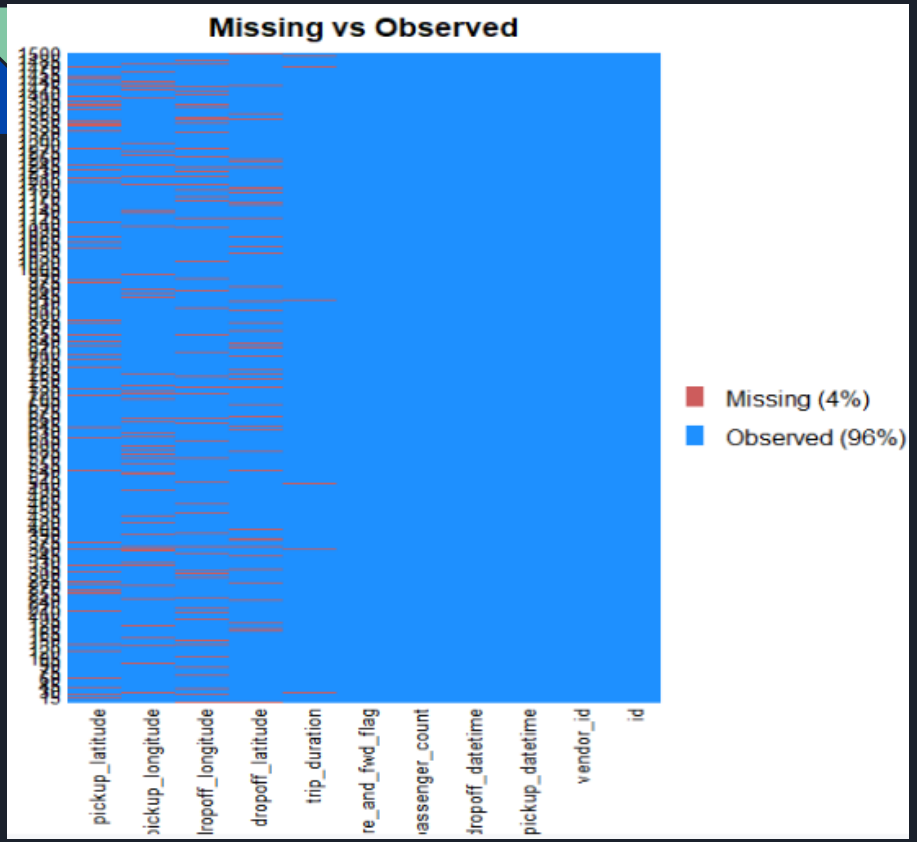

Missing Data Pattern

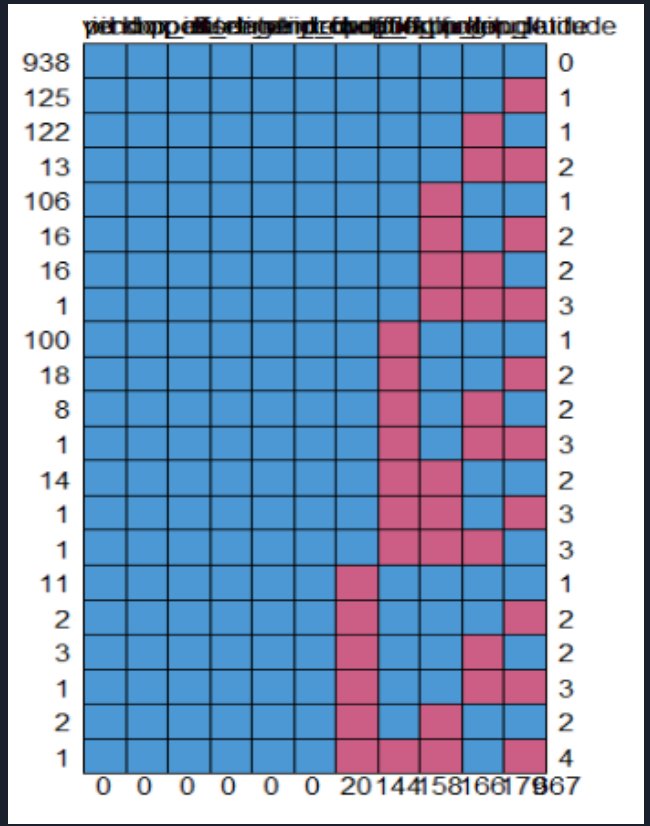

Data Acquisition | **Data Cleaning** | EDA | Building Data Models | Challenges to OLS | Model Validation | Prediction

Summary- After MICE

From summary, we can see that NAs are removed after multiple imputation.

```
      id            vendor_id pickup_datetime
 Length:1500        1:719     Min.   :2016-01-01 00:20:00
 Class :character   2:781     1st Qu.:2016-02-19 19:27:45
 Mode  :character             Median :2016-04-03 20:27:30
                              Mean   :2016-04-02 15:30:06
                              3rd Qu.:2016-05-14 20:18:15
                              Max.   :2016-06-30 23:21:00
 dropoff_datetime              passenger_count pickup_longitude
 Min.   :2016-01-01 00:24:00   1:1044          Min.   :-74.19
 1st Qu.:2016-02-19 19:45:00   2: 239          1st Qu.:-73.99
 Median :2016-04-03 20:43:00   3:  64          Median :-73.98
 Mean   :2016-04-02 15:44:59   4:  32          Mean   :-73.97
 3rd Qu.:2016-05-14 20:41:15   5:  80          3rd Qu.:-73.97
 Max.   :2016-06-30 23:37:00   6:  41          Max.   :-73.78
 pickup_latitude dropoff_longitude dropoff_latitude store_and_fwd_flag
 Min.   :40.63   Min.   :-74.18    Min.   :40.59    N:1492
 1st Qu.:40.74   1st Qu.:-73.99    1st Qu.:40.73    Y:   8
 Median :40.75   Median :-73.98    Median :40.76
 Mean   :40.75   Mean   :-73.97    Mean   :40.75
 3rd Qu.:40.77   3rd Qu.:-73.96    3rd Qu.:40.77
 Max.   :40.88   Max.   :-73.78    Max.   :41.00
 trip_duration
 Min.   :   78.0
 1st Qu.:  431.0
 Median :  680.0
 Mean   :  902.5
 3rd Qu.: 1073.5
 Max.   :86137.0
```

Feature Engineering

- **Neighborhoods Incorporation**:

  We have assigned a neighborhood to every pickup and dropoff location .

  In order to determine the neighborhoods of given locations, we have used the publicly available Zillow Neighborhood Boundary Shapefiles.

- **Distance Incorporation**:

  We have calculated distance_miles using the Longitude and Latitude.

Skewness of Target Variable- Trip Duration

- Trip duration has almost log normal distribution.
- Most of the rides were less than 17 minutes.
- One potential outlier which is making our distribution bimodal i.e distribution with two peaks.
- It has edge-peak distribution i.e. a peak towards the edge of the distribution.

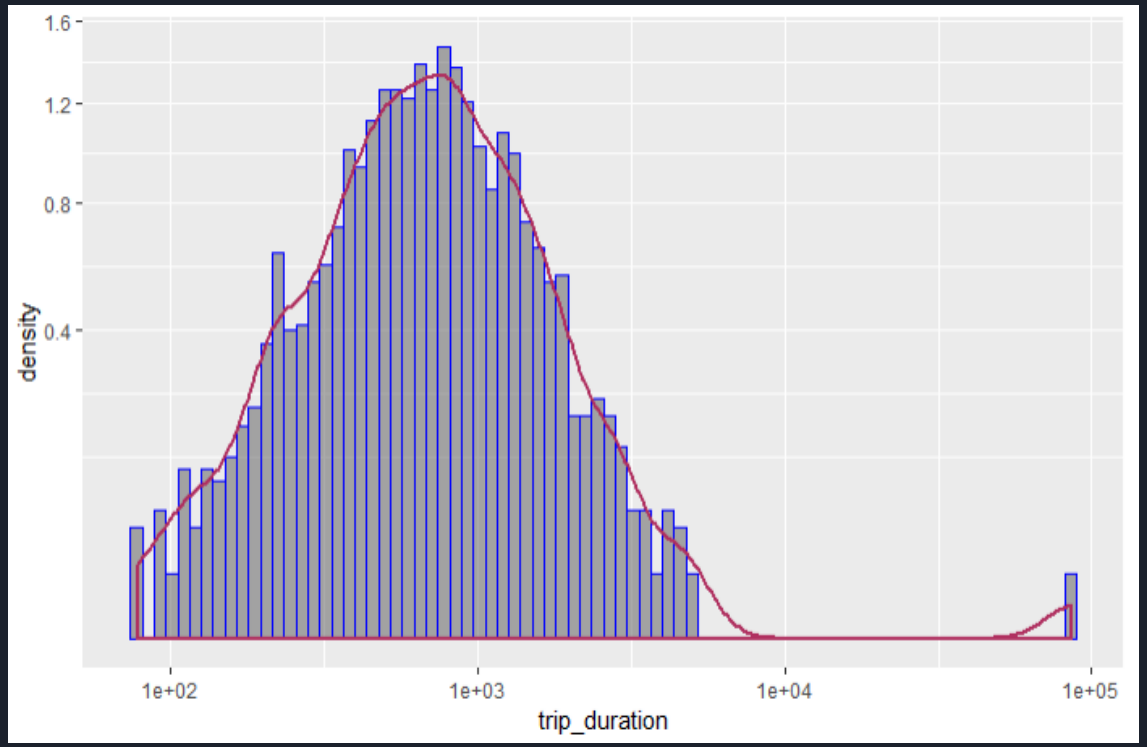

Box Plot of Trip Duration

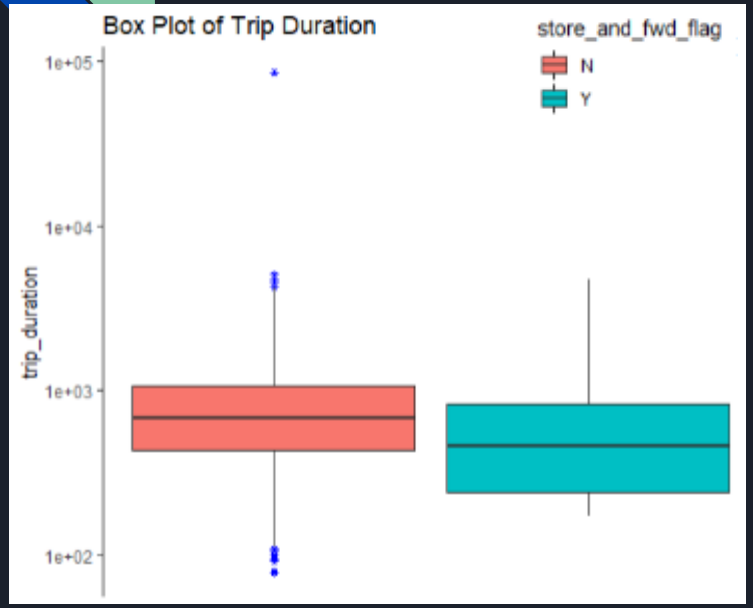

Normal QQ Plot of Trip Duration

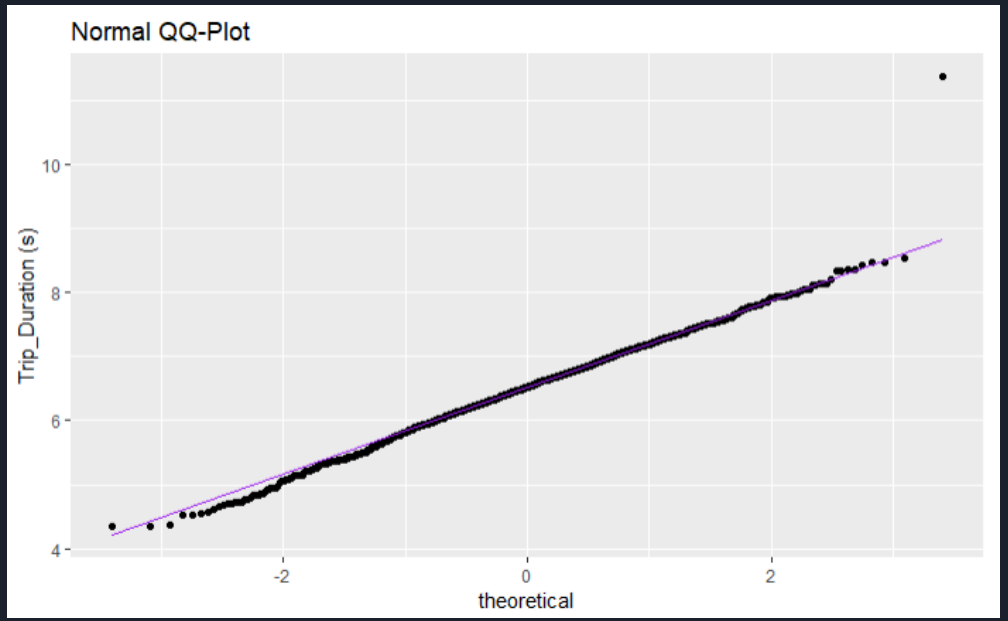

Skewness of Predictor Variables

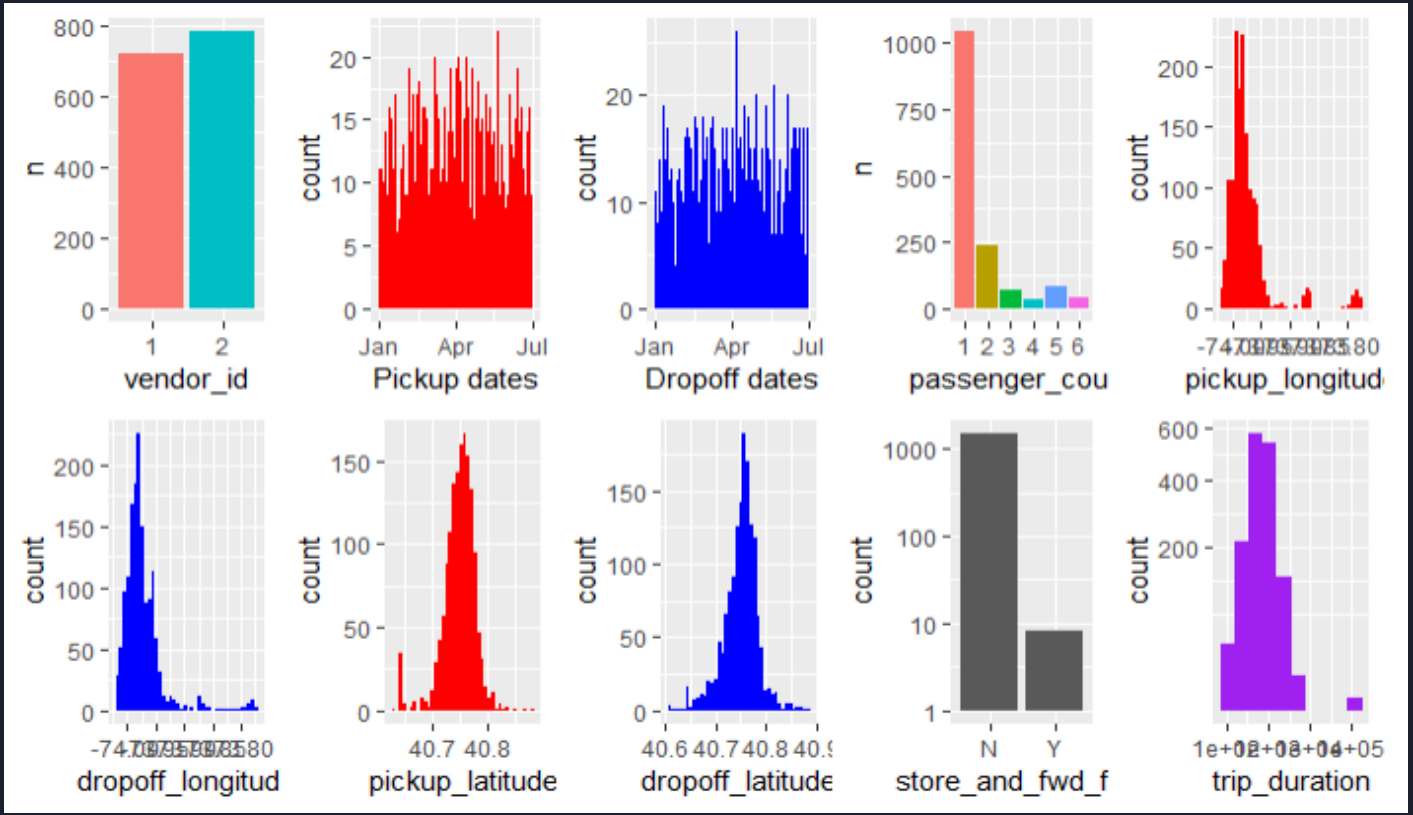

Skewness of Continuous Variables

- Except Trip duration, all the other variables are approximately normally distributed with few outliers.

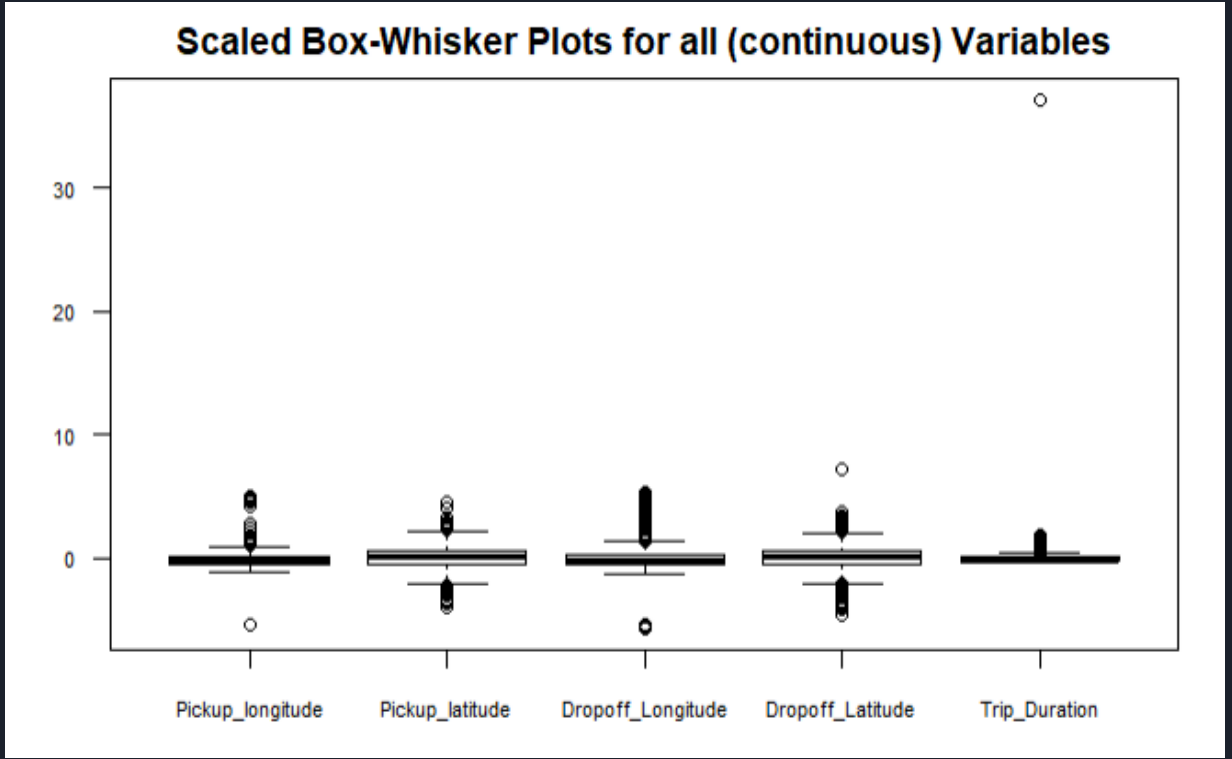

Distance (in Miles) Predictor

- Graph shows that Trip Duration has high correlation with Distance_miles.
- Due to the presence of an influential outlier, the graph is right skewed.

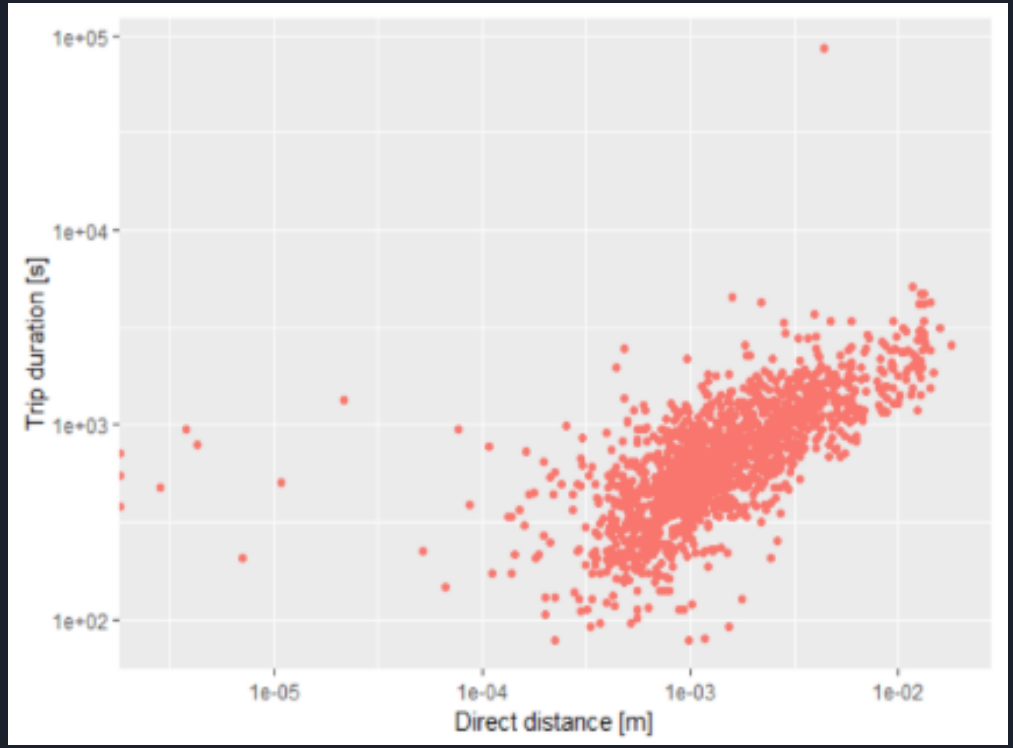

Heatmap based on Neighborhood

- The average trip duration is really low for rides within the neighborhood, it is logical.
- The historical average trip duration between two neighborhood certainly informs about the trip duration between those two neighborhoods in the future - Strong Predictor!

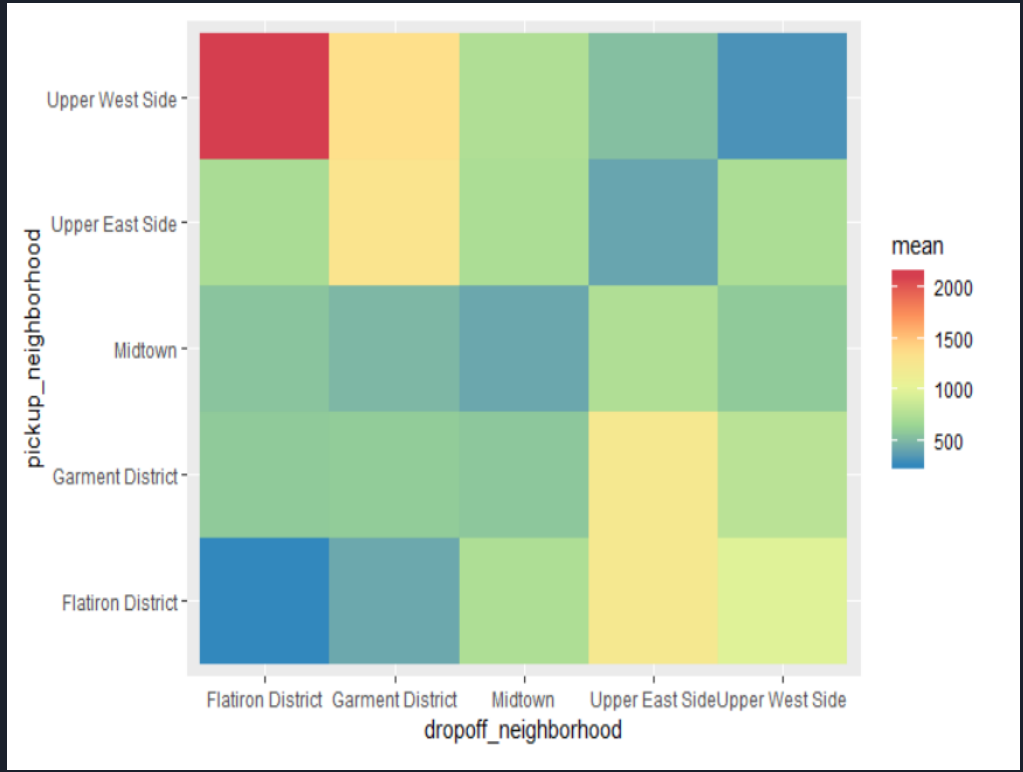

Leaflet

- A map of NYC and overlay a manageable number of pickup coordinates to get a general overview of the location and distances.

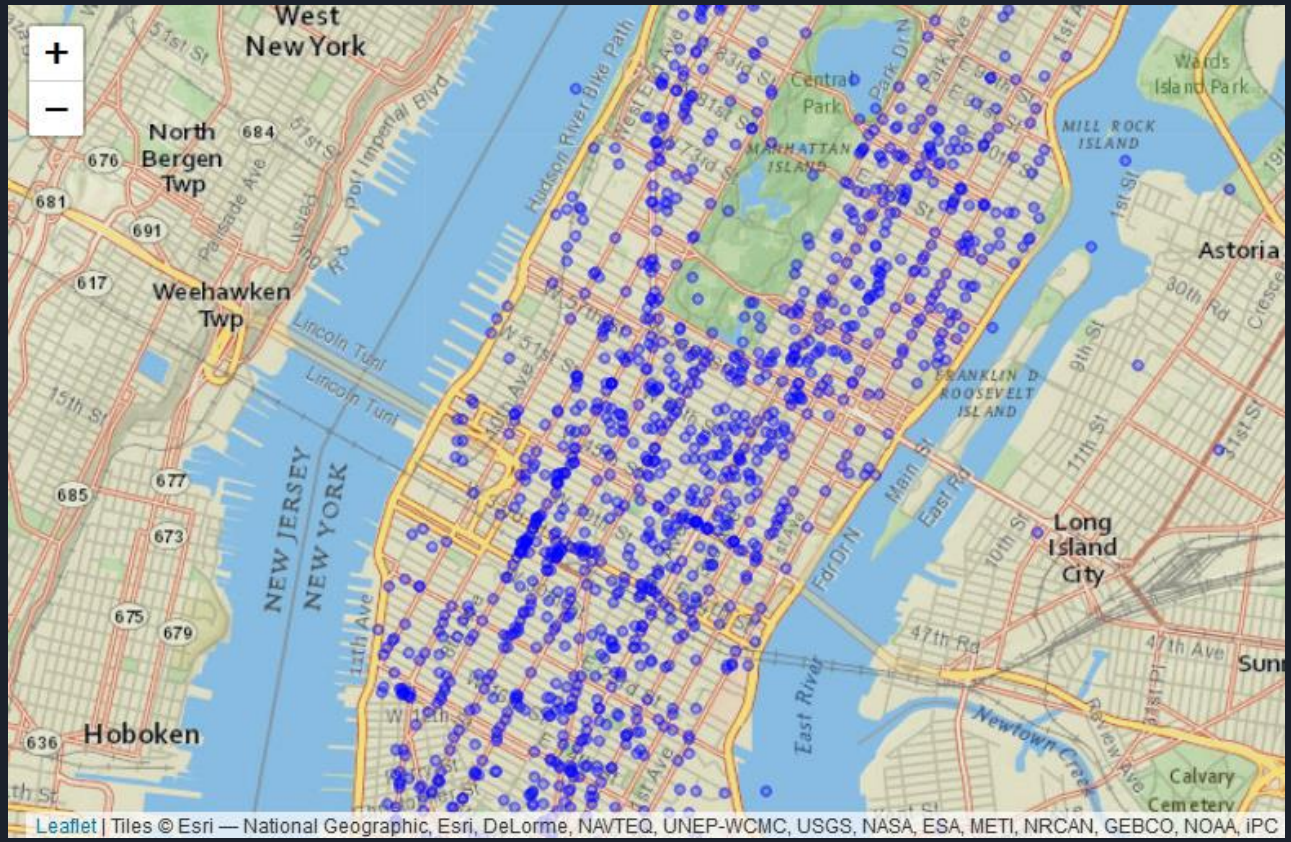

Analysis- Number of Rides

- Almost all the days of the weeks and hours of the day, vendor 2 has more pickup as compared to vendor 1.
- Friday is the busiest day.
- Monday has the lowest number of rides.

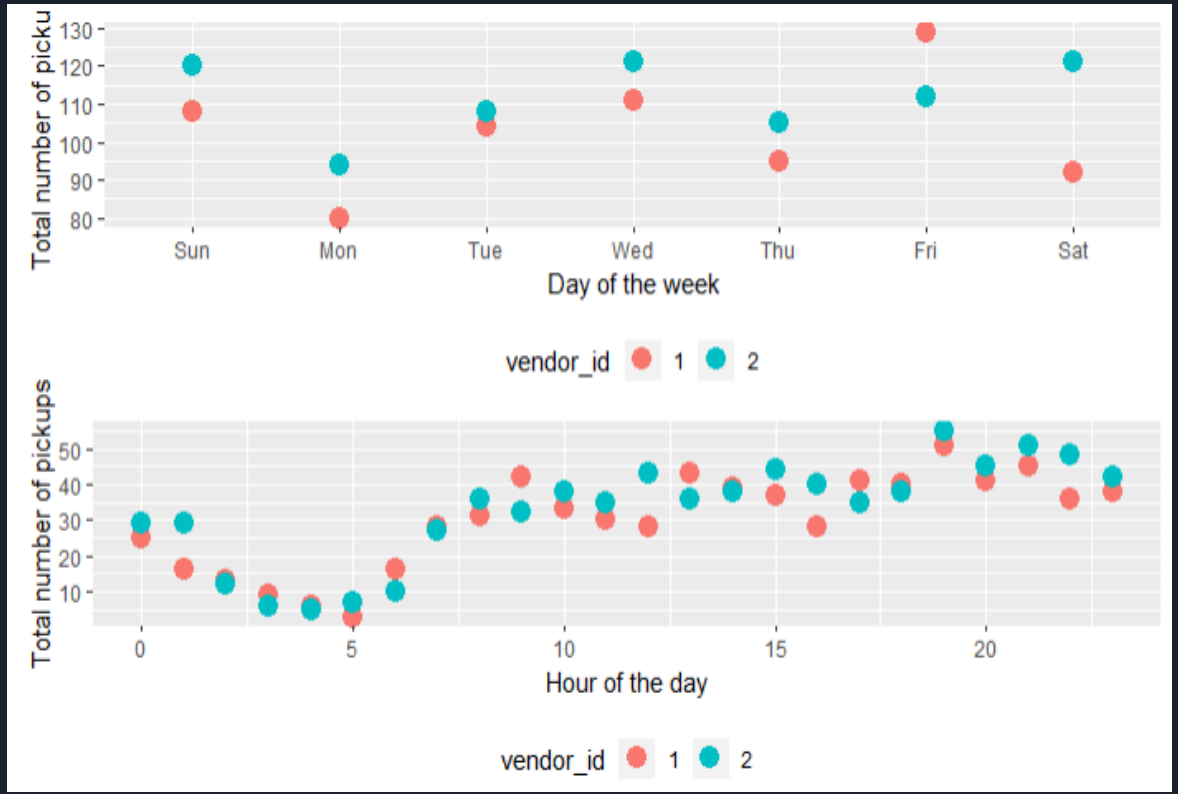

Analysis- Number of Rides based on Hour of the Day

- On friday, Saturday and Sunday, we have more trips during early morning hours on the contrary, trips are low in between 5 to 10.
- One possible reason for this distinction could be the contrast between the lifestyle of people on business days and night life.

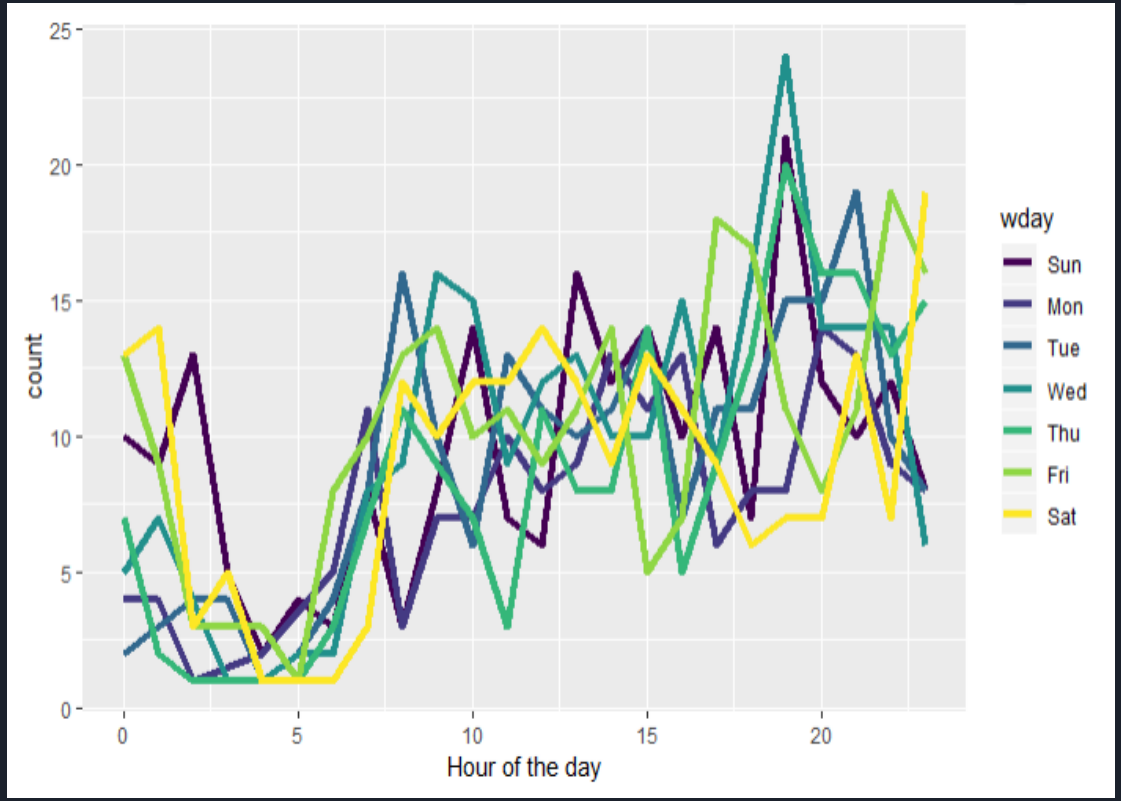

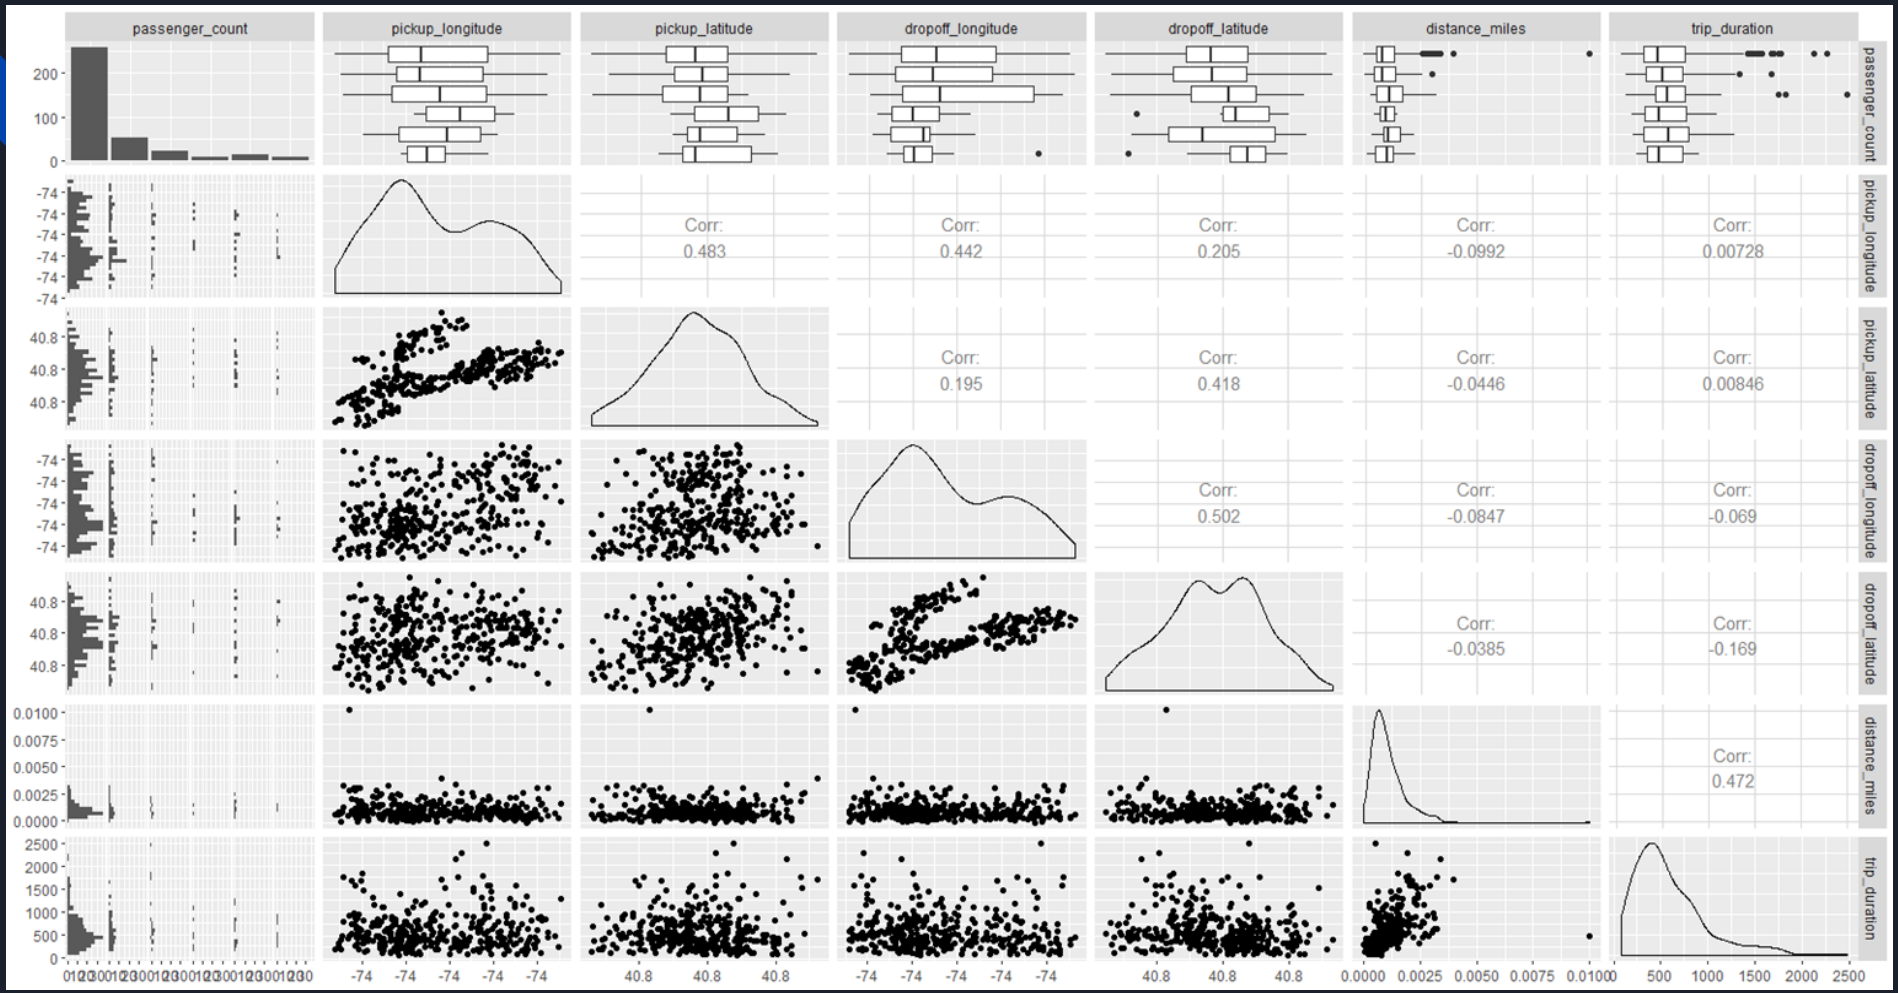
passenger_count
pickup_longitude
pickup_latitude
dropoff_longitude
dropoff_latitude
distance_miles
trip_duration
Corr: 0.483
Corr: 0.442
Corr: 0.205
Corr: -0.0992
Corr: 0.00728
Corr: 0.195
Corr: 0.418
Corr: -0.0446
Corr: 0.00846
Corr: 0.502
Corr: -0.0847
Corr: -0.069
Corr: -0.0385
Corr: -0.169
Corr: 0.472

GGpairs- Two-at-time Redundancy and Outliers

- No two predictors are redundant.
- Distance_miles is right skewed due to presence of an influential outlier.
- Pickup_Latitude and Dropoff_Latitude has approximately normal distribution as from the leaflet, we can see that most of the rides has their pickup location of Manhattan.
- Pickup_longitutde and Dropoff_longitutde has a right tailed distribution. This is consistent with the fact that most of the rides were for Manhattan.

Corrplot- Strong Predictor Checked

- Trip Duration has highest correlation with Distance_miles
- Although correlation of neighborhood (categorical variable) and Trip duration is not possible, the data reveals that there is correlation between the two.
- Strong Predictors - Neighborhood and iles

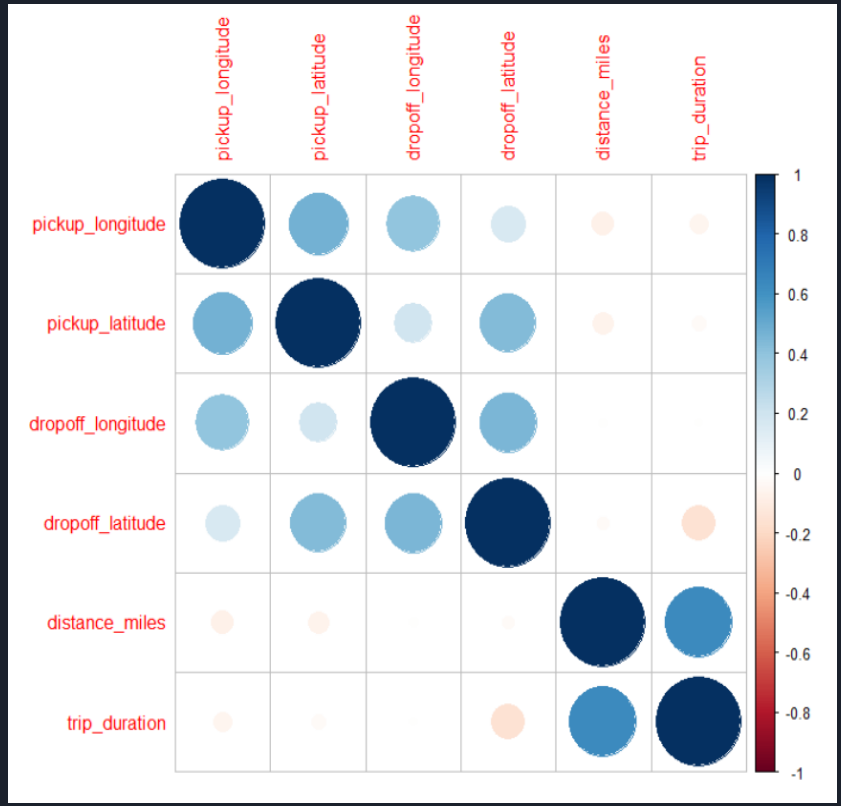

Model Fitting

- g1 : Fitted biggest model using Linear Regression

```
> summary(g1)

Call:
lm(formula = trip_duration ~ passenger_count + pickup_longitude +
    pickup_latitude + dropoff_longitude + dropoff_latitude +
    distance_miles + pickup_neighborhood + dropoff_neighborhood +
    same_neighborhood, data = TAXIDATANEW)

Residuals:
   Min     1Q Median     3Q    Max
 -1462   -192    -49    129   1823

Coefficients:
                                          Estimate Std. Error t value Pr(>|t|)
(Intercept)                             -3.75e+05   5.50e+05   -0.68   0.4955
passenger_count2                         5.45e+01   4.99e+01    1.09   0.2754
passenger_count3                         1.61e+02   7.28e+01    2.21   0.0274 *
passenger_count4                         3.99e+01   1.35e+02    0.29   0.7684
passenger_count5                        -8.50e+01   9.50e+01   -0.90   0.3713
passenger_count6                         7.54e+01   1.18e+02    0.64   0.5214
pickup_longitude                        -5.00e+03   3.87e+03   -1.29   0.1979
pickup_latitude                          9.61e+03   3.67e+03    2.62   0.0093 **
dropoff_longitude                        1.98e+03   3.87e+03    0.51   0.6085
dropoff_latitude                        -5.85e+03   3.54e+03   -1.65   0.0992 .
distance_miles                           1.54e+05   2.47e+04    6.23  1.4e-09 ***
pickup_neighborhoodGarment District      4.87e+00   9.20e+01    0.05   0.9578
pickup_neighborhoodMidtown              -1.63e+01   8.89e+01   -0.18   0.8546
pickup_neighborhoodUpper East Side       7.59e+01   1.27e+02    0.60   0.5497
pickup_neighborhoodUpper West Side      -1.62e+02   1.51e+02   -1.07   0.2837
dropoff_neighborhoodGarment District     2.30e+02   9.16e+01    2.51   0.0126 *
dropoff_neighborhoodMidtown              1.38e+02   8.59e+01    1.60   0.1095
dropoff_neighborhoodUpper East Side      1.49e+02   1.27e+02    1.18   0.2394
dropoff_neighborhoodUpper West Side      1.14e+02   1.48e+02    0.77   0.4414
same_neighborhood1                      -2.29e+02   4.39e+01   -5.22  3.2e-07 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 322 on 339 degrees of freedom
Multiple R-squared:  0.359,     Adjusted R-squared:  0.323
F-statistic:   10 on 19 and 339 DF,  p-value: <2e-16
```

Fit Plot and Residual Plot

- g1 is not a very good model as points are scattered.
- Residual plot shows a strong pattern

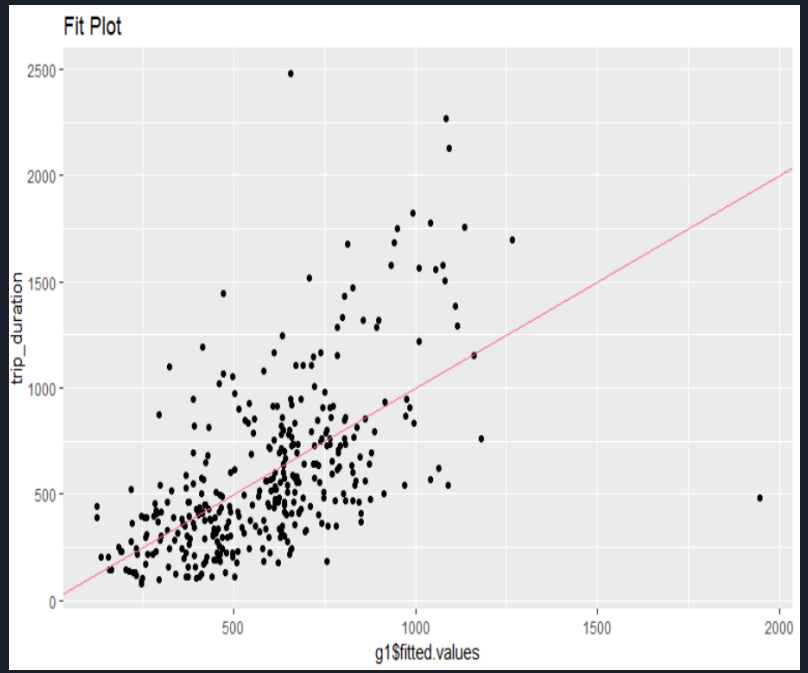

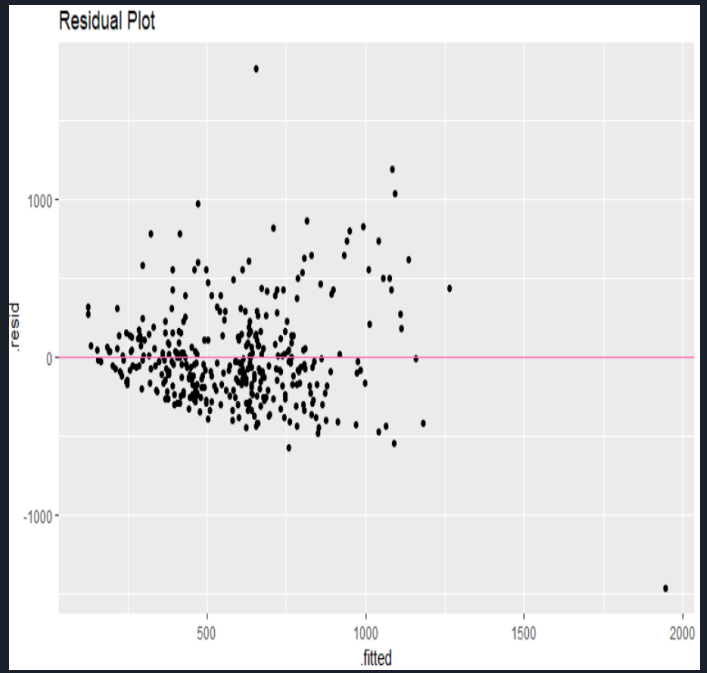

Model Fitting

- g2 : Applied Stepwise Regression on the biggest model (g1).

```
Step:  AIC=4157
trip_duration ~ pickup_latitude + dropoff_latitude + distance_miles +
    dropoff_neighborhood + same_neighborhood

                     Df Sum of Sq      RSS  AIC
<none>                             36454515 4157
- dropoff_latitude      1    359672 36814187 4158
- dropoff_neighborhood  4   1600896 38055411 4164
- pickup_latitude       1   1080815 37535330 4165
- same_neighborhood     1   2944757 39399273 4183
- distance_miles        1   4646674 41101189 4198
```

```
> summary(g2)

Call:
lm(formula = trip_duration ~ pickup_latitude + dropoff_latitude +
    distance_miles + dropoff_neighborhood + same_neighborhood,
    data = TAXIDATANEW)

Residuals:
    Min      1Q  Median      3Q     Max
-1488.4  -190.9   -46.8   115.1  2040.9

Coefficients:
                                        Estimate Std. Error t value Pr(>|t|)
(Intercept)                              32487.0   129053.5    0.25   0.8014
pickup_latitude                           4696.3     1457.9    3.22   0.0014 **
dropoff_latitude                         -5484.3     2951.3   -1.86   0.0640 .
distance_miles                          160948.7    24096.7    6.68  9.5e-11 ***
dropoff_neighborhoodGarment District       218.5       86.0    2.54   0.0115 *
dropoff_neighborhoodMidtown                163.4       83.0    1.97   0.0498 *
dropoff_neighborhoodUpper East Side        228.3      107.9    2.12   0.0350 *
dropoff_neighborhoodUpper West Side        119.2      138.1    0.86   0.3884
same_neighborhood1                        -227.1       42.7   -5.32  1.9e-07 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 323 on 350 degrees of freedom
Multiple R-squared:  0.336,     Adjusted R-squared:  0.32
F-statistic: 22.1 on 8 and 350 DF,  p-value: <2e-16
```

Model Fitting

- g5 : Applied Stepwise Regression with one predictor variable as log transformed.

```
> summary(g5)

Call:
lm(formula = trip_duration ~ pickup_latitude + dropoff_latitude +
    dropoff_neighborhood + log(distance_miles), data = TAXIDATANEW,
    na.action = na.exclude)

Residuals:
   Min     1Q Median     3Q    Max
-950.7 -205.5  -69.9  115.3 2016.9

Coefficients:
                                      Estimate Std. Error t value Pr(>|t|)
(Intercept)                            83355.4   127247.9    0.66    0.513
pickup_latitude                         3621.0     1450.4    2.50    0.013 *
dropoff_latitude                       -5607.3     2935.8   -1.91    0.057 .
dropoff_neighborhoodGarment District     286.7       86.4    3.32    0.001 **
dropoff_neighborhoodMidtown              166.5       82.3    2.02    0.044 *
dropoff_neighborhoodUpper East Side      164.1      107.8    1.52    0.129
dropoff_neighborhoodUpper West Side      102.4      137.6    0.74    0.457
log(distance_miles)                      273.3       23.7   11.54   <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 325 on 351 degrees of freedom
Multiple R-squared:  0.323,     Adjusted R-squared:  0.31
F-statistic: 23.9 on 7 and 351 DF,  p-value: <2e-16
```

Model Fitting

- g6 : Applied Stepwise regression with target variable as log transformed.

```
> summary(g6)

Call:
lm(formula = log(trip_duration) ~ pickup_latitude + dropoff_latitude +
    pickup_neighborhood + dropoff_neighborhood + distance_miles,
    data = TAXIDATANEW)

Residuals:
   Min     1Q Median     3Q    Max
-3.666 -0.305  0.049  0.352  1.892

Coefficients:
                                  Estimate Std. Error t value Pr(>|t|)
(Intercept)                       238.0027   291.0824    0.82   0.4141
pickup_latitude                     7.8334     5.4039    1.45   0.1481
dropoff_latitude                  -13.5395     5.0626   -2.67   0.0078 **
pickup_neighborhoodGarment District   0.1021     0.1508    0.68   0.4990
pickup_neighborhoodMidtown            0.1210     0.1520    0.80   0.4268
pickup_neighborhoodUpper East Side    0.0111     0.1921    0.06   0.9540
pickup_neighborhoodUpper West Side   -0.2270     0.2507   -0.91   0.3657
dropoff_neighborhoodGarment District   0.3648     0.1481    2.46   0.0143 *
dropoff_neighborhoodMidtown           0.4738     0.1438    3.29   0.0011 **
dropoff_neighborhoodUpper East Side   0.3867     0.1914    2.02   0.0441 *
dropoff_neighborhoodUpper West Side   0.4474     0.2386    1.88   0.0616 .
distance_miles                    392.1135    37.4987   10.46   <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.554 on 347 degrees of freedom
Multiple R-squared:  0.319,     Adjusted R-squared:  0.297
F-statistic: 14.8 on 11 and 347 DF,  p-value: <2e-16
```

Model Fitting

- g7 : Applied Stepwise regression on the big model with target variable as log transformed.

```
> summary(g7)

Call:
lm(formula = log(trip_duration) ~ passenger_count + pickup_longitude +
    pickup_latitude + dropoff_longitude + dropoff_latitude +
    distance_miles + pickup_neighborhood + dropoff_neighborhood +
    same_neighborhood, data = TAXIDATANEW)

Residuals:
    Min      1Q  Median      3Q     Max
-2.1341 -0.3299  0.0269  0.3101  1.7706

Coefficients:
                                        Estimate Std. Error t value Pr(>|t|)
(Intercept)                          -387.1072   891.0788   -0.43   0.6643
passenger_count2                        0.1561     0.0808    1.93   0.0542 .
passenger_count3                        0.2128     0.1179    1.80   0.0720 .
passenger_count4                        0.1181     0.2193    0.54   0.5906
passenger_count5                       -0.0943     0.1538   -0.61   0.5403
passenger_count6                        0.2472     0.1904    1.30   0.1949
pickup_longitude                       -2.1366     6.2714   -0.34   0.7335
pickup_latitude                         9.0422     5.9474    1.52   0.1294
dropoff_longitude                      -3.1036     6.2696   -0.50   0.6209
dropoff_latitude                       -8.9152     5.7349   -1.55   0.1210
distance_miles                        262.7512    40.0801    6.56   2.1e-10 ***
pickup_neighborhoodGarment District     0.0730     0.1490    0.49   0.6247
pickup_neighborhoodMidtown              0.0897     0.1440    0.62   0.5335
pickup_neighborhoodUpper East Side      0.0979     0.2052    0.48   0.6338
pickup_neighborhoodUpper West Side     -0.1558     0.2445   -0.64   0.5244
dropoff_neighborhoodGarment District    0.2908     0.1483    1.96   0.0507 .
dropoff_neighborhoodMidtown             0.3688     0.1391    2.65   0.0084 **
dropoff_neighborhoodUpper East Side     0.4628     0.2053    2.25   0.0248 *
dropoff_neighborhoodUpper West Side     0.3064     0.2393    1.28   0.2013
same_neighborhood1                     -0.4796     0.0711   -6.75   6.4e-11 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.522 on 339 degrees of freedom
Multiple R-squared:  0.41,      Adjusted R-squared:  0.376
F-statistic: 12.4 on 19 and 339 DF,  p-value: <2e-16
```

Compare Coefficients

- Comparing the coefficients, we realize that model 4 i.e. g6 is better for predicting Trip_duration

```
> compareCoefs(g1,g2,g5,g6,g7,se=FALSE)
Calls:
1: lm(formula = trip_duration ~ passenger_count + pickup_longitude + pickup_latitude + dropoff_longitude
  + dropoff_latitude + distance_miles + pickup_neighborhood + dropoff_neighborhood + same_neighborhood,
  data = TAXIDATANEW)
2: lm(formula = trip_duration ~ pickup_latitude + dropoff_latitude + distance_miles +
  dropoff_neighborhood + same_neighborhood, data = TAXIDATANEW)
3: lm(formula = trip_duration ~ pickup_longitude + pickup_latitude + dropoff_latitude +
  pickup_neighborhood + dropoff_neighborhood + log(distance_miles), data = TAXIDATANEW, na.action =
  na.exclude)
4: lm(formula = log(trip_duration) ~ pickup_latitude + dropoff_latitude + pickup_neighborhood +
  dropoff_neighborhood + distance_miles, data = TAXIDATANEW)
5: lm(formula = log(trip_duration) ~ passenger_count + pickup_longitude + pickup_latitude +
  dropoff_longitude + dropoff_latitude + distance_miles + pickup_neighborhood + dropoff_neighborhood +
  same_neighborhood, data = TAXIDATANEW)
```

	Model 1	Model 2	Model 3	Model 4	Model 5
(Intercept)	-375462	32487	-577900	238	-387
passenger_count2	54.496				0.156
passenger_count3	161.265				0.213
passenger_count4	39.898				0.118
passenger_count5	-85.0160				-0.0943
passenger_count6	75.439				0.247
pickup_longitude	-4996.25		-4607.52		-2.14
pickup_latitude	9608.82	4696.32	10799.71	7.83	9.04
dropoff_longitude	1984.5				-3.1
dropoff_latitude	-5854.17	-5484.35	-4925.88	-13.54	-8.92
distance_miles	154142	160949		392	263
pickup_neighborhoodGarment District	4.874		44.494	0.102	0.073
pickup_neighborhoodMidtown	-16.3071		-54.3136	0.1210	0.0897
pickup_neighborhoodUpper East Side	75.8879		-11.7016	0.0111	0.0979
pickup_neighborhoodUpper West Side	-162.070		-254.838	-0.227	-0.156
dropoff_neighborhoodGarment District	229.594	218.504	268.257	0.365	0.291
dropoff_neighborhoodMidtown	137.883	163.420	146.204	0.474	0.369
dropoff_neighborhoodUpper East Side	149.437	228.339	125.154	0.387	0.463
dropoff_neighborhoodUpper West Side	113.876	119.220	72.052	0.447	0.306
same_neighborhood1	-228.88	-227.06			-0.48
log(distance_miles)			262		

Partial F-test - Transformed Target Variable

- Anova of g6 and g7
- As P- Value is less than the alpha (0.05), we conclude that the model g6 is better

```
> anova(g6,g7)
Analysis of Variance Table

Model 1: log(trip_duration) ~ pickup_latitude + dropoff_latitude + pickup_neighborhood +
    dropoff_neighborhood + distance_miles
Model 2: log(trip_duration) ~ passenger_count + pickup_longitude + pickup_latitude +
    dropoff_longitude + dropoff_latitude + distance_miles + pickup_neighborhood +
    dropoff_neighborhood + same_neighborhood
  Res.Df   RSS Df Sum of Sq   F  Pr(>F)
1    347 106.4
2    339  92.2  8      14.2 6.5 6.9e-08 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1
```

Partial F-test - Transformed Predictor Variable

- Anova of g5 and g1
- As P- Value is greater than the alpha (0.05), we conclude that the smaller model, with one predictor as log transformed, is better i.e. g5 is better

```
> anova(g5,g1)
Analysis of Variance Table

Model 1: trip_duration ~ pickup_longitude + pickup_latitude + dropoff_latitude +
    pickup_neighborhood + dropoff_neighborhood + log(distance_miles)
Model 2: trip_duration ~ passenger_count + pickup_longitude + pickup_latitude +
    dropoff_longitude + dropoff_latitude + distance_miles + pickup_neighborhood +
    dropoff_neighborhood + same_neighborhood
  Res.Df      RSS Df Sum of Sq    F Pr(>F)
1    346 36238191
2    339 35161708  7   1076483 1.48   0.17
```

Non-Constant Variance

The t-test does not reject constant error variance with a level of significance 5%, since the p-value, 0.404, is greater than 0.05

```
> summary(lm(abs(residuals(g7)) ~ fitted(g7)))

Call:
lm(formula = abs(residuals(g7)) ~ fitted(g7))

Residuals:
    Min      1Q  Median      3Q     Max
-0.3936 -0.2625 -0.0686  0.1891  1.8151

Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.5994     0.2495    2.40    0.017 *
fitted(g7)    -0.0337     0.0403   -0.84    0.404
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.323 on 357 degrees of freedom
Multiple R-squared:  0.00195,   Adjusted R-squared:  -0.000844
F-statistic: 0.698 on 1 and 357 DF,  p-value: 0.404
```

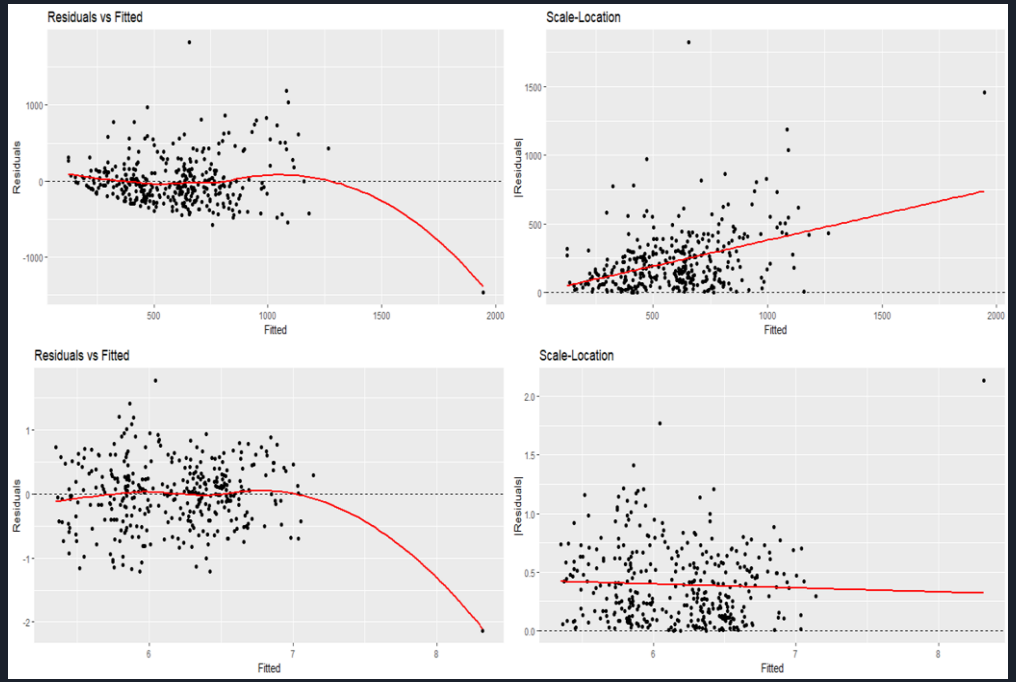

Non-Normal Errors

Normal QQ-plots for detecting

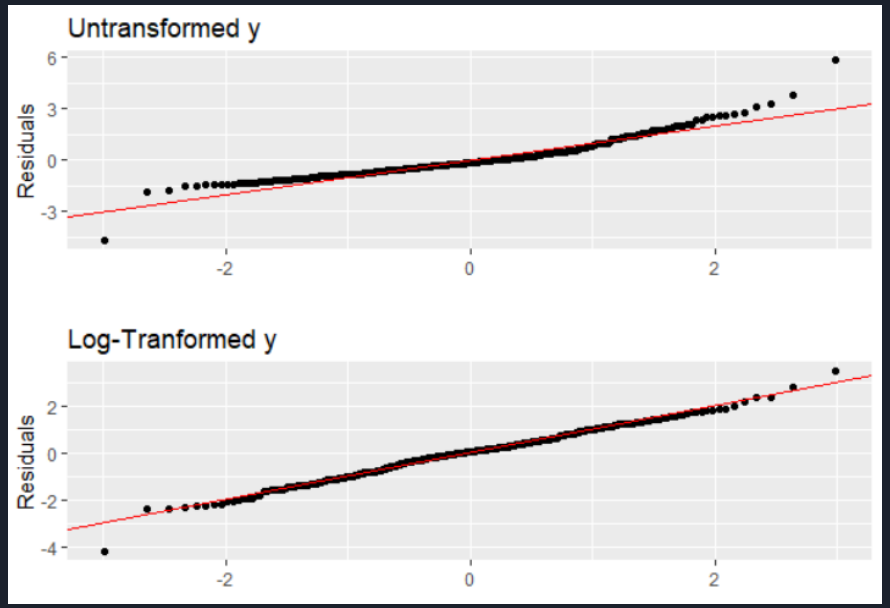


```
> shapiro.test(residuals(g1))

        Shapiro-Wilk normality test

data:  residuals(g1)
W = 0.9, p-value = 1e-13

> shapiro.test(residuals(g7))

        Shapiro-Wilk normality test

data:  residuals(g7)
W = 1, p-value = 0.06
```

We fail to reject the null hypothesis of normality for the residuals of log-transformed model with level of significance 5% since the p-value is greater than 0.05.

Non-Normal Errors

Box-Cox Power Transformation

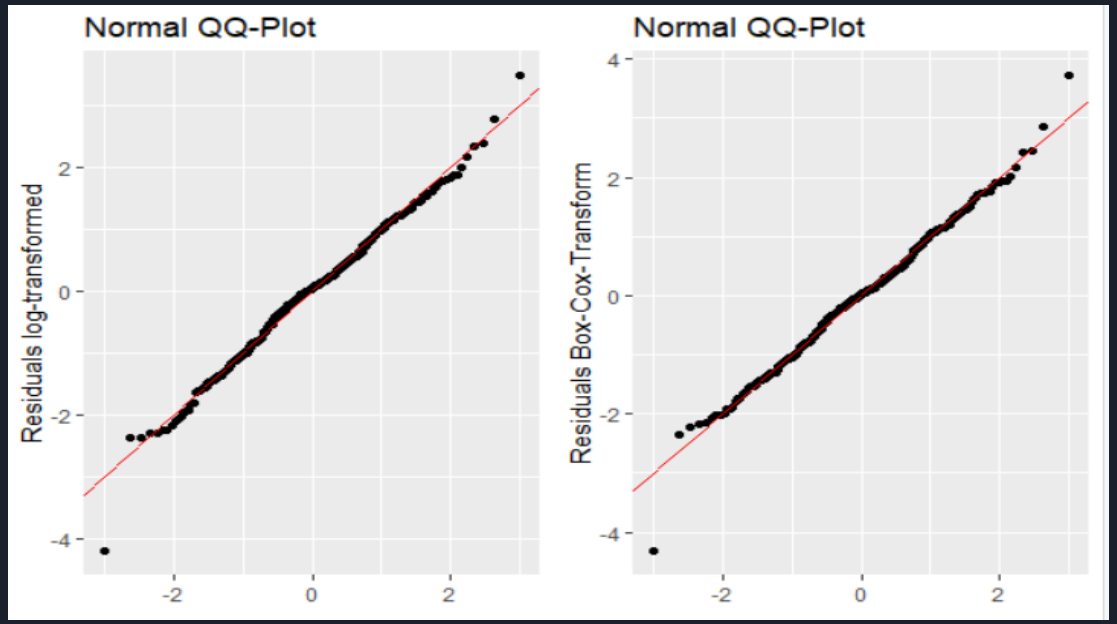


```
> shapiro.test(residuals(glam))

        Shapiro-Wilk normality test

data:  residuals(glam)
W = 1, p-value = 0.05
```

The Shapiro-Wilk test concludes that the errors are normal for the Box-Cox Transformed model with level of significance 5% since the p-value is approximately equal to 0.05.

Influential Outliers

Influence Plot

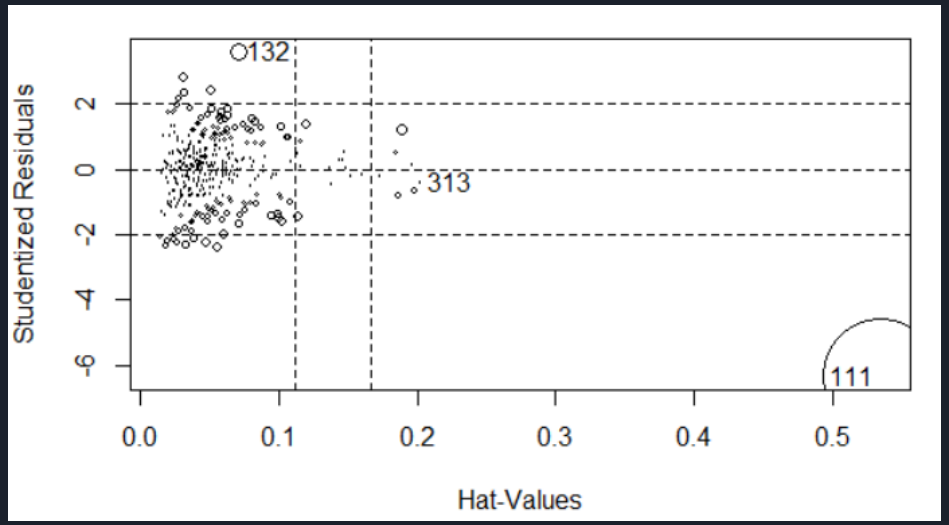

Half Normal Plot

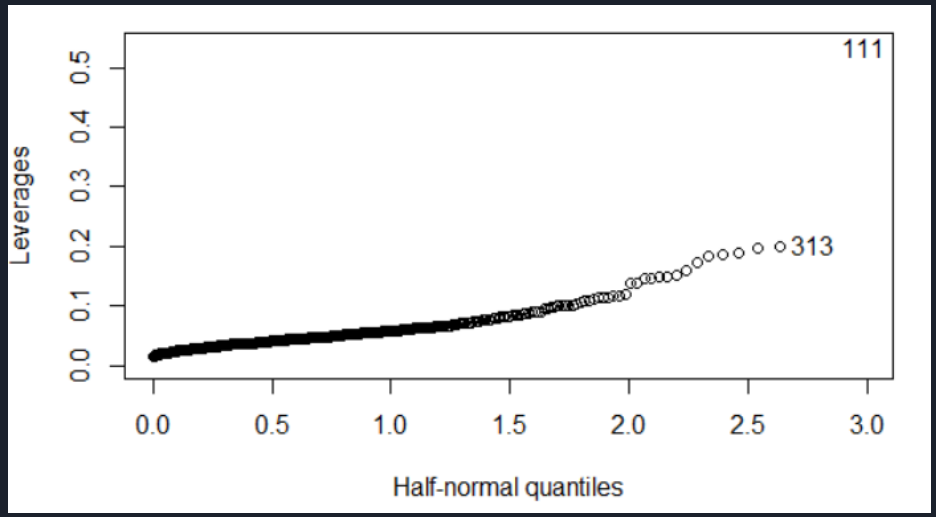

Influential Outlier

Half normal plot of Cook's Distance

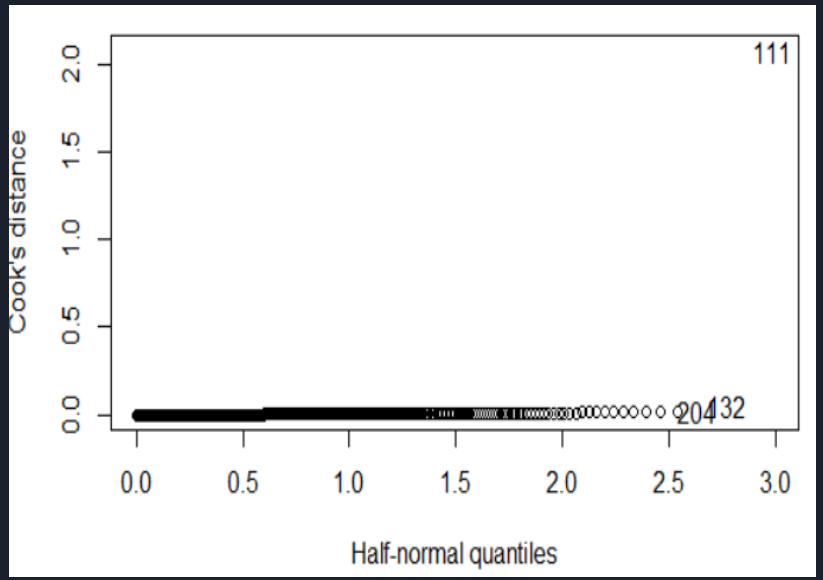

Influence Plot (without 111 data point)

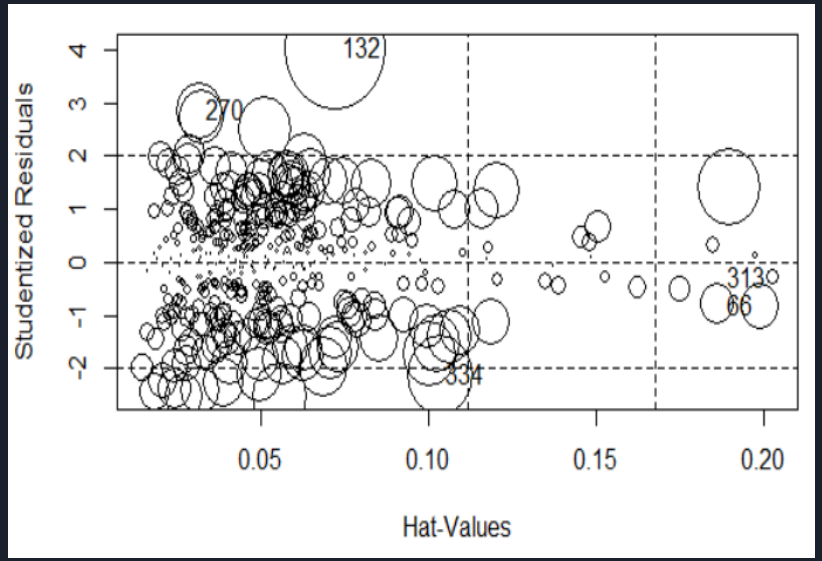

Omnibus diagnostic plot function

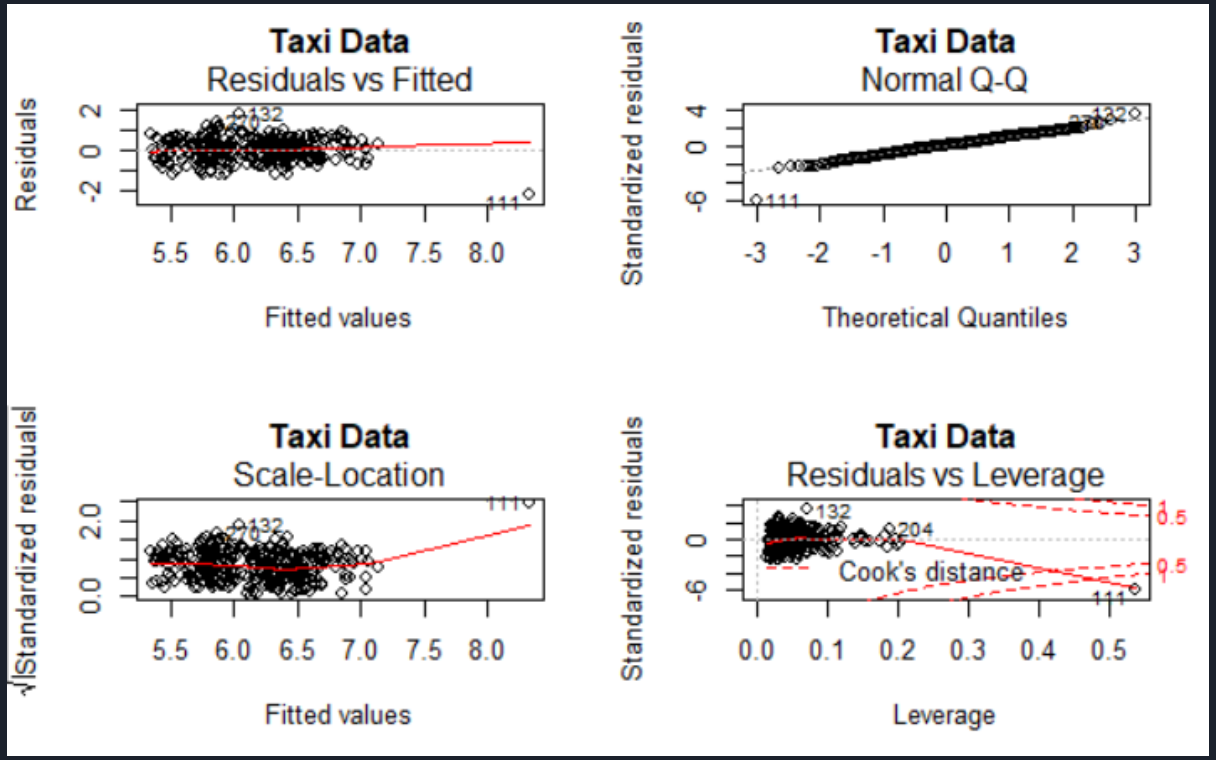

Correct Model Specification

Target Variable: Trip Duration

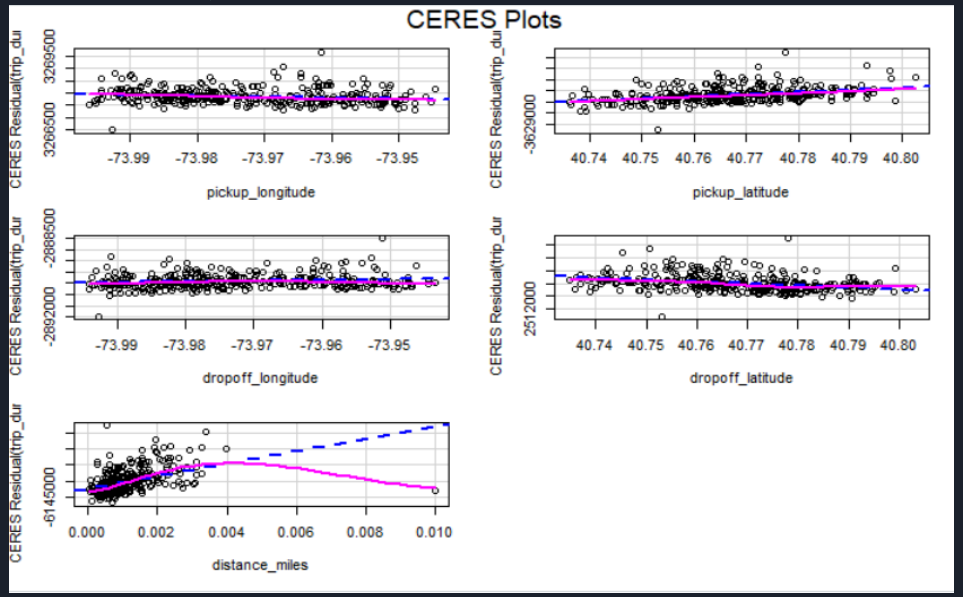

Target Variable: log(Trip Duration)

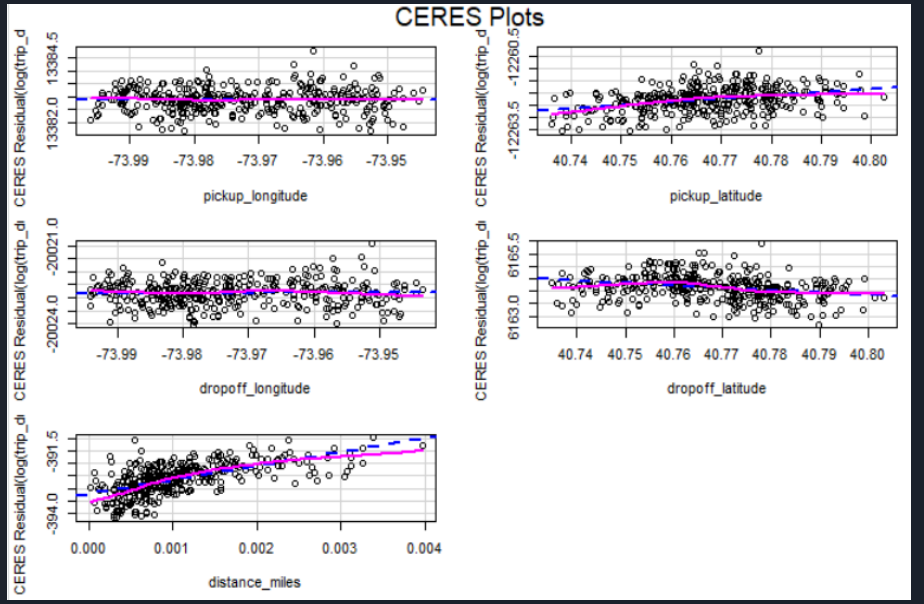

After log transformation, the model looks normal

Cross Validation of linear models

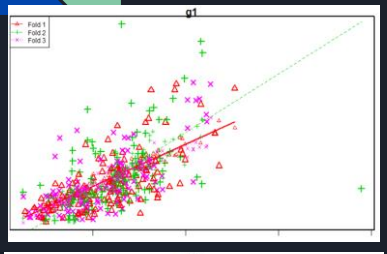

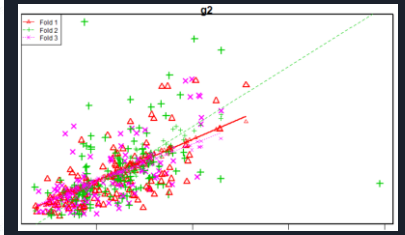

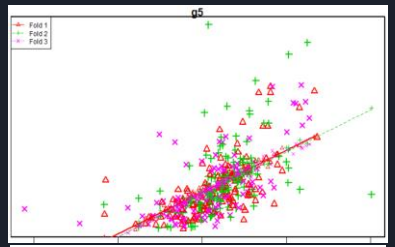

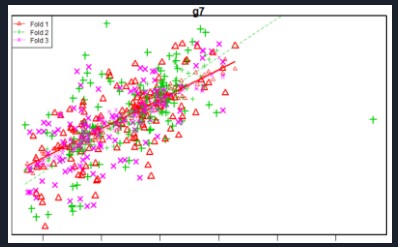

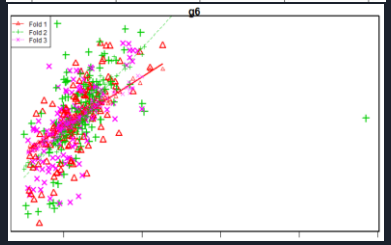

	mse.g1	mse.g2	mse.g5	mse.g6	mse.g7
1	135711	132367	112996	0.339	0.365
2	130854	125073	110115	0.328	0.364
3	136415	128714	112871	0.357	0.395
4	148878	141865	114177	0.358	0.384
5	128400	123075	105864	0.333	0.367
6	134255	125216	109379	0.331	0.364
7	123855	118842	114576	0.323	0.375
8	135330	132155	111722	0.331	0.364
9	128358	124701	108729	0.326	0.376
10	133711	126080	113462	0.357	0.398

From the cross validation predicted values, we can see that the model g6 holds better mse overall.

*Rest of the graphs is in Rmd File.

Individual Confidence Interval

- We are focusing in the 95% CI on the model g1 .

```
> confint(g1)
                                                  2.5 %  97.5 %
(Intercept)                                   -1.46e+06  706804
passenger_count2                              -4.36e+01     153
passenger_count3                               1.80e+01     304
passenger_count4                              -2.26e+02     306
passenger_count5                              -2.72e+02     102
passenger_count6                              -1.56e+02     307
pickup_longitude                              -1.26e+04    2621
pickup_latitude                                2.39e+03   16832
dropoff_longitude                             -5.63e+03    9599
dropoff_latitude                              -1.28e+04    1111
distance_miles                                 1.05e+05  202822
pickup_neighborhoodGarment District           -1.76e+02     186
pickup_neighborhoodMidtown                    -1.91e+02     159
pickup_neighborhoodUpper East Side            -1.73e+02     325
pickup_neighborhoodUpper West Side            -4.59e+02     135
dropoff_neighborhoodGarment District           4.95e+01     410
dropoff_neighborhoodMidtown                   -3.11e+01     307
dropoff_neighborhoodUpper East Side           -1.00e+02     399
dropoff_neighborhoodUpper West Side           -1.77e+02     404
same_neighborhood1                            -3.15e+02    -143
```

Comparing CI of Coefficients with and without Bonferroni Models

- Comparing the model g and model g where Trip duration is converted to factor.

```
> confint(g9, level =.95)
                      2.5 %    97.5 %
(Intercept)       16388511 17245884
pickup_longitude    121483    127935
pickup_latitude    -149169   -141489
dropoff_latitude    -41802    -40095
distance_miles     1088651   1131783
> confint(g9, level  = 1-0.05/(2*6))
                     0.208 % 99.792 %
(Intercept)       16188559 17445836
pickup_longitude    119979    129439
pickup_latitude    -150960   -139698
dropoff_latitude    -42200    -39696
distance_miles     1078592   1141842
> confint(g10, level=.95)
                      2.5 %    97.5 %
(Intercept)       16388511 17245884
pickup_longitude    121483    127935
pickup_latitude    -149169   -141489
dropoff_latitude    -41802    -40095
distance_miles     1088651   1131783
> confint(g10, level = 1-0.05/(2*6))
                     0.208 % 99.792 %
(Intercept)       16188559 17445836
pickup_longitude    119979    129439
pickup_latitude    -150960   -139698
dropoff_latitude    -42200    -39696
distance_miles     1078592   1141842
```

Thanks

Any Questions?